

## ✓ Introduction to Linear Regression

### Module 1, Lab 1: Understanding Linear Regression

Linear regression is one of the most fundamental algorithms in machine learning. In this lab, you will learn how to implement and understand linear regression from the ground up.

#### Learning Objectives

- Understand the mathematical foundation of linear regression
- Implement linear regression using scikit-learn
- Evaluate model performance using appropriate metrics
- Visualize the relationship between features and target variables

#### Business Problem

We will predict house prices based on various features like size, location, and number of bedrooms. This is a classic regression problem that helps understand how different factors influence property values.

## ✓ Setup

### Installing and Importing Libraries

First, we need to install the necessary libraries for our analysis.

```
# Install required packages
!pip install --upgrade pip
!pip install pandas numpy matplotlib seaborn scikit-learn
```

```
# Import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
from sklearn.preprocessing import StandardScaler
import warnings
warnings.filterwarnings('ignore')
```

```
# Set plotting style
plt.style.use('seaborn-v0_8')
%matplotlib inline
```

## ✓ Loading and Exploring the Data

We'll use the Boston Housing dataset, which is a classic dataset for regression problems.

```
# Create a synthetic housing dataset since Boston Housing has ethical concerns
np.random.seed(42)
n_samples = 500

# Generate synthetic features
size = np.random.normal(2500, 1000, n_samples) # House size in sq ft
bedrooms = np.random.randint(1, 3, n_samples) # Number of bedrooms
age = np.random.randint(1, 40, n_samples) # Age of house
location_score = np.random.uniform(1, 15, n_samples) # Location desirability score



# Generate target variable (price) with realistic relationships
price = (size * 150 + bedrooms * 10000 - age * 500 + location_score * 5000 +
         np.random.normal(0, 20000, n_samples))

# Create DataFrame
df = pd.DataFrame({
    'size': size,
    'bedrooms': bedrooms,
    'age': age,
    'location_score': location_score,
    'price': price
})

# Ensure positive prices
df['price'] = np.maximum(df['price'], 50000)

print(f"Dataset shape: {df.shape}")
df.head()
```

Dataset shape: (500, 5)

	size	bedrooms	age	location_score	price	
0	2996.714153	2	2	7.732654	526140.555227	
1	2361.735699	2	3	1.570557	381310.199779	
2	3147.688538	2	23	3.295993	483461.368381	
3	4023.029856	1	12	7.140501	657347.655162	
4	2265.846625	2	20	11.147572	401388.131272	

Next steps:

[Generate code with df](#)[New interactive sheet](#)

## ✓ Exploratory Data Analysis

Let's explore our dataset to understand the relationships between variables.

```
# Basic statistics
print("Dataset Info:")
print(df.info())
print("\nBasic Statistics:")
print(df.describe())
```

```
Dataset Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   size                   500 non-null   float64
1   bedrooms               500 non-null   int64
2   age                    500 non-null   int64
3   location_score         500 non-null   float64
4   price                  500 non-null   float64
dtypes: float64(3), int64(2)
memory usage: 19.7 KB
None
```

Basic Statistics:

	size	bedrooms	age	location_score	price
count	500.000000	500.0000	500.000000	500.000000	500.000000
mean	2506.837995	1.4800	20.706000	7.842993	419596.027891
std	981.253247	0.5001	11.057309	3.833236	148733.192245
min	-741.267340	1.0000	1.000000	1.004655	50000.000000
25%	1799.692596	1.0000	11.000000	4.521573	314973.790697
50%	2512.797146	1.0000	21.000000	7.763655	417765.925236
75%	3136.783254	2.0000	30.000000	10.730839	518168.632731
max	6352.731491	2.0000	39.000000	14.936727	960282.865982

```
# Visualize the relationships
fig, axes = plt.subplots(2, 2, figsize=(15, 10))


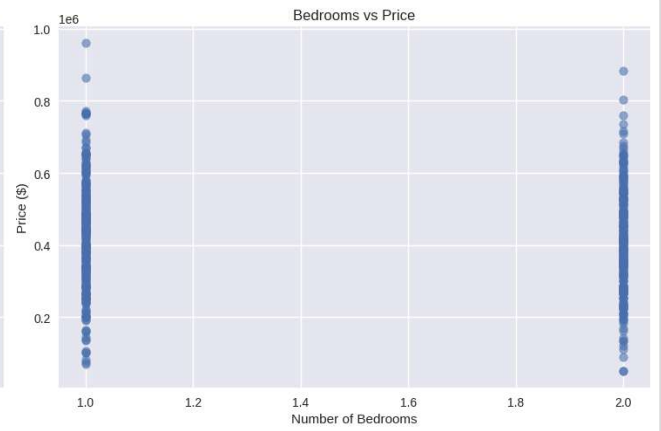
# Size vs Price
axes[0, 0].scatter(df['size'], df['price'], alpha=0.6)
axes[0, 0].set_xlabel('House Size (sq ft)')
axes[0, 0].set_ylabel('Price ($)')
axes[0, 0].set_title('House Size vs Price')

# Bedrooms vs Price
axes[0, 1].scatter(df['bedrooms'], df['price'], alpha=0.6)
axes[0, 1].set_xlabel('Number of Bedrooms')
axes[0, 1].set_ylabel('Price ($)')
axes[0, 1].set_title('Bedrooms vs Price')

# Age vs Price
axes[1, 0].scatter(df['age'], df['price'], alpha=0.6)
axes[1, 0].set_xlabel('House Age (years)')
axes[1, 0].set_ylabel('Price ($)')
axes[1, 0].set_title('House Age vs Price')

# Location Score vs Price
axes[1, 1].scatter(df['location_score'], df['price'], alpha=0.6)
axes[1, 1].set_xlabel('Location Score')
axes[1, 1].set_ylabel('Price ($)')
axes[1, 1].set_title('Location Score vs Price')

plt.tight_layout()
plt.show()
```





## Building the Linear Regression Model

Now let's build our linear regression model step by step.

```
# Prepare features and target
X = df[['size', 'bedrooms', 'age', 'location_score']]
y = df['price']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_sta

print(f"Training set size: {X_train.shape[0]}")
print(f"Test set size: {X_test.shape[0]}")
```

```
Training set size: 400
Test set size: 100
```

```
# Create and train the linear regression model
model = LinearRegression()
model.fit(X_train, y_train)

# Make predictions
y_pred_train = model.predict(X_train)
y_pred_test = model.predict(X_test)

print("Model trained successfully!")
```

```
Model trained successfully!
```

## Model Evaluation

Let's evaluate how well our model performs.

```
# Calculate evaluation metrics
train_mse = mean_squared_error(y_train, y_pred_train)
test_mse = mean_squared_error(y_test, y_pred_test)
train_r2 = r2_score(y_train, y_pred_train)
test_r2 = r2_score(y_test, y_pred_test)
train_mae = mean_absolute_error(y_train, y_pred_train)
test_mae = mean_absolute_error(y_test, y_pred_test)

print("Model Performance:")
print(f"Training R2 Score: {train_r2:.4f}")
print(f"Test R2 Score: {test_r2:.4f}")
print(f"Training RMSE: ${np.sqrt(train_mse):.2f}")
print(f"Test RMSE: ${np.sqrt(test_mse):.2f}")
```

```
print(f"Training MAE: ${train_mae:,.2f}")
print(f"Test MAE: ${test_mae:,.2f}")
```


Model Performance:  
Training R<sup>2</sup> Score: 0.9842  
Test R<sup>2</sup> Score: 0.9745  
Training RMSE: \$18,826.43  
Test RMSE: \$22,962.57  
Training MAE: \$14,772.90  
Test MAE: \$17,011.32

```
# Visualize predictions vs actual values
plt.figure(figsize=(12, 5))

# Training set
plt.subplot(1, 2, 1)
plt.scatter(y_train, y_pred_train, alpha=0.6)
plt.plot([y_train.min(), y_train.max()], [y_train.min(), y_train.max()], 'r--', lw=
plt.xlabel('Actual Price')
plt.ylabel('Predicted Price')
plt.title(f'Training Set: Actual vs Predicted\nR2 = {train_r2:.4f}')

# Test set
plt.subplot(1, 2, 2)
plt.scatter(y_test, y_pred_test, alpha=0.6)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--', lw=2)
plt.xlabel('Actual Price')
plt.ylabel('Predicted Price')
plt.title(f'Test Set: Actual vs Predicted\nR2 = {test_r2:.4f}')

plt.tight_layout()
plt.show()
```



## ✓ Understanding the Model Coefficients

Let's examine what our model learned about the relationship between features and house prices.

```
# Display model coefficients
feature_names = X.columns
coefficients = model.coef_
intercept = model.intercept_

print("Linear Regression Equation:")
print(f"Price = {intercept:.2f}", end="")
for name, coef in zip(feature_names, coefficients):
    print(f" + ({coef:.2f} × {name})", end="")
print("\n")

# Create a DataFrame for better visualization
coef_df = pd.DataFrame({
    'Feature': feature_names,
    'Coefficient': coefficients,
    'Abs_Coefficient': np.abs(coefficients)
```



```
}).sort_values('Abs_Coefficient', ascending=False)

print("Feature Importance (by coefficient magnitude):")
print(coef_df)
```

Linear Regression Equation:

Price = -5364.60 + (150.86 × size) + (10051.77 × bedrooms) + (-497.18 × age) + (5209

Feature Importance (by coefficient magnitude):

	Feature	Coefficient	Abs_Coefficient
1	bedrooms	10051.769439	10051.769439
3	location_score	5209.693806	5209.693806
2	age	-497.178564	497.178564
0	size	150.855713	150.855713

```
# Visualize feature importance
plt.figure(figsize=(10, 6))
colors = ['green' if x > 0 else 'red' for x in coefficients]
plt.barh(feature_names, coefficients, color=colors, alpha=0.7)
plt.xlabel('Coefficient Value')
plt.title('Feature Coefficients in Linear Regression Model')
plt.axvline(x=0, color='black', linestyle='-', alpha=0.3)
plt.grid(True, alpha=0.3)
plt.show()
```

## Feature Coefficients in Linear Regression Model

## ✓ Challenge: Make Predictions

Now it's your turn! Use the trained model to make predictions for new houses.

```
# Challenge: Predict prices for these houses
new_houses = pd.DataFrame({
    'size': [1800, 2500, 1200],
    'bedrooms': [3, 4, 2],
    'age': [10, 5, 25],
    'location_score': [8.5, 9.2, 6.0]
})

# TODO: Use the model to predict prices for these houses
# predictions = model.predict(new_houses)

# Uncomment the lines below after making predictions
# for i, price in enumerate(predictions):
#     print(f"House {i+1}: ${price:,.2f}")

print("Houses to predict:")
print(new_houses)
```